

Floating-Point Representation

CS3010

Numerical Methods

Dr. Amar Raheja

Section 1.3

Lecture 3

Number System: Bases

- Most Popular Number Systems:
- Decimal: Base 10 (Digits of number system?)
 - 0-9
- Octal: Base 8 (Digits of number system?)
 - 0-7
- Binary: Base 2 (Digits of number system?)
 - 0-1
- Hexadecimal: Base 16 (Digits of number system?)
 - (0-9, A-F)

Number Representation Examples

- Other Bases to Decimal
- Base 10: $4586_{10} = 6 \times 10^0 + 8 \times 10^1 + 6 \times 10^2 + 4 \times 10^3$
- Base 8: $56327_8 = 7 \times 8^0 + 2 \times 8^1 + 3 \times 8^2 + 6 \times 8^3 + 5 \times 8^4 = 23767$
- Base 2: $(10111)_2 = 1 \times 2^0 + 1 \times 2^1 + 1 \times 2^2 + 0 \times 2^3 + 1 \times 2^4 = 23_{10}$
- Base 16: $(4A59F)_{16} = 15 \times 16^0 + 9 \times 16^1 + 5 \times 16^2 + 10 \times 16^3 + 4 \times 16^4 = 304543$
- Decimal to Bases
- Decimal to Binary: Divide by 2, put remainder in a stack (right most bit-least significant), divide quotient by 2 and keep doing this till quotient is 0
- Other conversions are similar and division is done by appropriate base

Convert Base 10 Integer to binary representation

Table 1 Converting a base-10 integer to binary representation.

| | Quotient | Remainder |
|------|----------|-----------|
| 11/2 | 5 | $1 = a_0$ |
| 5/2 | 2 | $1 = a_1$ |
| 2/2 | 1 | $0 = a_2$ |
| 1/2 | 0 | $1 = a_3$ |

Hence
$$\begin{aligned}(11)_{10} &= (a_3 a_2 a_1 a_0)_2 \\ &= (1011)_2\end{aligned}$$

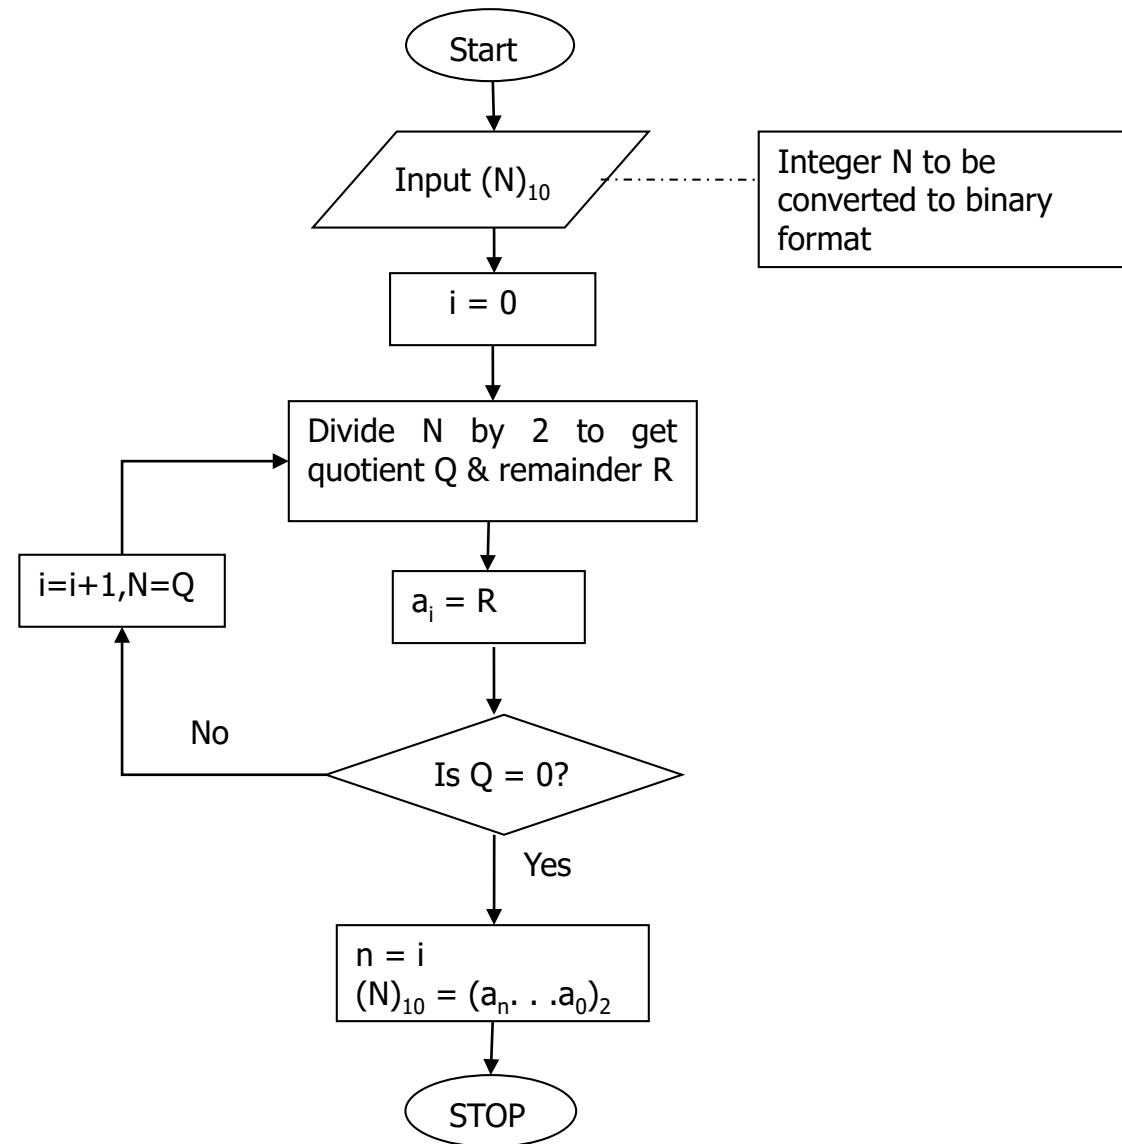
Convert Base 10 Integer to Octal representation

Table 1 Converting a base-10 integer to binary representation.

| | Quotient | Remainder |
|-------|----------|-----------|
| 145/8 | 18 | 1 |
| 18/8 | 2 | 2 |
| 2/8 | 0 | 2 |

Hence, $(145)_{10} = (221)_8$

Flowchart



Fractional Decimal Number to Binary

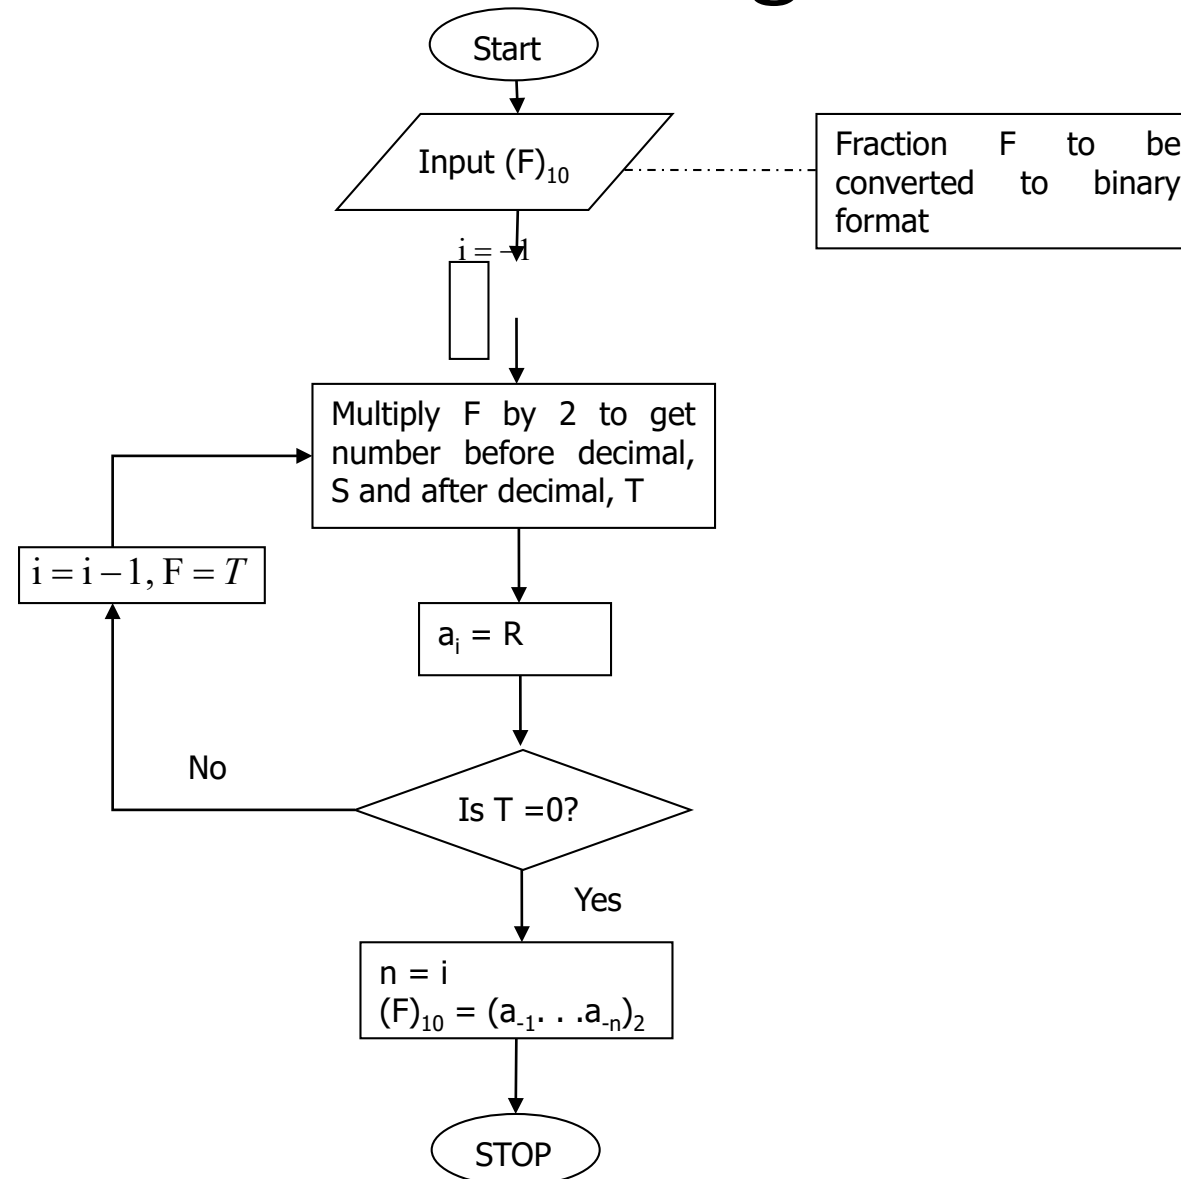
Table 2. Converting a base-10 fraction to binary representation.

| | Number | Number after decimal | Number before decimal |
|-------------------|--------|----------------------|-----------------------|
| 0.1875×2 | 0.375 | 0.375 | $0 = a_{-1}$ |
| 0.375×2 | 0.75 | 0.75 | $0 = a_{-2}$ |
| 0.75×2 | 1.5 | 0.5 | $1 = a_{-3}$ |
| 0.5×2 | 1.0 | 0.0 | $1 = a_{-4}$ |

Hence

$$\begin{aligned}(0.1875)_{10} &= (a_{-1}a_{-2}a_{-3}a_{-4})_2 \\ &= (0.0011)_2\end{aligned}$$

Flowchart for converting Fractional Part



Decimal Number to Binary

$$(11.1875)_{10} = (?.)_2$$

Since $(11)_{10} = (1011)_2$

and $(0.1875)_{10} = (0.0011)_2$

then, one can combine these two to get

$$(11.1875)_{10} = (1011.0011)_2$$

All Fractional Decimal Numbers Cannot be Represented Exactly

- 0.3 or 0.1 cannot be represented in a finite way in binary

Table 3. Converting a base-10 fraction to approximate binary representation.

| | Number | Number after decimal | Number before Decimal |
|----------------|--------|----------------------|-----------------------|
| 0.3×2 | 0.6 | 0.6 | $0 = a_{-1}$ |
| 0.6×2 | 1.2 | 0.2 | $1 = a_{-2}$ |
| 0.2×2 | 0.4 | 0.4 | $0 = a_{-3}$ |
| 0.4×2 | 0.8 | 0.8 | $0 = a_{-4}$ |
| 0.8×2 | 1.6 | 0.6 | $1 = a_{-5}$ |

$$(0.3)_{10} \approx (a_{-1}a_{-2}a_{-3}a_{-4}a_{-5})_2 = (0.01001)_2 = 0.28125$$

Another Way to Look at Conversion

- Convert $(11.1875)_{10}$ to base 2

$$\begin{aligned}(11)_{10} &= 2^3 + 3 \\&= 2^3 + 2^1 + 1 \\&= 2^3 + 2^1 + 2^0 \\&= 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 \\&= (1011)_2\end{aligned}$$

$$\begin{aligned}(0.1875)_{10} &= 2^{-3} + 0.0625 \\&= 2^{-3} + 2^{-4} \\&= 0 \times 2^{-1} + 0 \times 2^{-2} + 1 \times 2^{-3} + 1 \times 2^{-4} \\&= (0.0011)_2\end{aligned}$$

Scientific Notation

- Scientific Notation: Shifting the decimal point to represent decimal numbers

256.78 is written as $+ 2.5678 \times 10^2$

0.003678 is written as $+ 3.678 \times 10^{-3}$

-256.78 is written as $- 2.5678 \times 10^2$

- Normalized Scientific Notation: The number is represented by a fraction multiplied by 10^n , and the leading digit in the fraction is not zero (except when the number involved is zero).
- One can write 79325 as 0.79325×10^5 , not 0.07932×10^6 or 7.9325×10^4 or any other way.

Normalized Scientific Notation

$$37.21829 = 0.37218\ 29 \times 10^2$$

$$0.00227\ 1828 = 0.22718\ 28 \times 10^{-2}$$

$$-30\ 00527.11059 = -0.30005\ 27110\ 59 \times 10^7$$

- So, change number to become a fractional number which is between 0 and 1 and is called the mantissa and the exponent is the power of 10

- The form of $x = \pm 0.d_1 d_2 d_3 \dots \times 10^n = \pm r \times 10^n$

where $d_1 \neq 0$ and n is an integer (positive, negative, or zero). The numbers d_1, d_2, \dots are the decimal digits 0, 1, 2, 3, 4, 5, 6, 7, 8, and 9.

$$x = \text{sign} \times \text{mantissa} \times 10^n$$

- a sign that is either + or −, a number r in the interval $\left[\frac{1}{10}, 1\right)$ and an integer power of 10. The number r is called the normalized mantissa and n the exponent.

Normalized Floating-Point Representation for Binary

- The floating-point representation in the binary system is similar to that in the decimal system in several ways. If $x \neq 0$, it can be written as

$$x = \pm q \times 2^m \quad \left(\frac{1}{2} \leq q < 1 \right)$$

- The mantissa q would be expressed as a sequence of zeros or ones in the form $q = \pm 0.b_1b_2b_3 \dots \times 2^m$ where $b_1 \neq 0$. Hence, where $b_1 = 1$
- Every computer has only a finite word length and a finite total capacity, so only numbers with a finite number of digits can be represented.
- A number is allotted only one word of storage in the single-precision mode (two or more words in double or extended precision).
- Clearly, irrational numbers cannot be represented, nor can those rational numbers that do not fit the finite format imposed by the computer.

Machine Numbers

- The real numbers that are representable in a computer are called its **machine numbers**.
- Since any number used in calculations with a computer system must conform to the format of numbers in that system, it must have a finite expansion. Numbers that have a nonterminating expansion cannot be accommodated precisely.
- Moreover, a number that has a terminating expansion in one base may have a nonterminating expansion in another.
 - Another good example of this is the following simple fraction as given in the introductory example
 - $(0.1)_{10} = (0.06314\ 6314\ 6314\ 6314\ \dots)_8$
 $= (0.0\ 0011\ 0011\ 0011\ 0011\ 0011\ 0011\ 0011\ \dots)_2$

5 bit Example


- A floating-point numbers must be of the form $q = \pm(0.b_1b_2b_3)_2 \times 2^{\pm k}$, where b_1, b_2, b_3 , and m are allowed to have only the value 0 or 1.
- There are two choices for the \pm , two choices for b_1 , two choices for b_2 , two choices for b_3 , and three choices for the exponent.
- Thus, at first, one would expect $2 \times 2 \times 2 \times 2 \times 3 = 48$ different numbers.

Possible positive numbers represented

- However, there is some duplication. For example, the nonnegative numbers in this system are as follows:
- Altogether there are 31 distinct numbers in the system. The positive numbers obtained are shown on a line

| | | |
|----------------------------------|----------------------------------|--------------------------------------|
| $0.000 \times 2^0 = 0$ | $0.000 \times 2^1 = 0$ | $0.000 \times 2^{-1} = 0$ |
| $0.001 \times 2^0 = \frac{1}{8}$ | $0.001 \times 2^1 = \frac{1}{4}$ | $0.001 \times 2^{-1} = \frac{1}{16}$ |
| $0.010 \times 2^0 = \frac{2}{8}$ | $0.010 \times 2^1 = \frac{2}{4}$ | $0.010 \times 2^{-1} = \frac{2}{16}$ |
| $0.011 \times 2^0 = \frac{3}{8}$ | $0.011 \times 2^1 = \frac{3}{4}$ | $0.011 \times 2^{-1} = \frac{3}{16}$ |
| $0.100 \times 2^0 = \frac{4}{8}$ | $0.100 \times 2^1 = \frac{4}{4}$ | $0.100 \times 2^{-1} = \frac{4}{16}$ |
| $0.101 \times 2^0 = \frac{5}{8}$ | $0.101 \times 2^1 = \frac{5}{4}$ | $0.101 \times 2^{-1} = \frac{5}{16}$ |
| $0.110 \times 2^0 = \frac{6}{8}$ | $0.110 \times 2^1 = \frac{6}{4}$ | $0.110 \times 2^{-1} = \frac{6}{16}$ |
| $0.111 \times 2^0 = \frac{7}{8}$ | $0.111 \times 2^1 = \frac{7}{4}$ | $0.111 \times 2^{-1} = \frac{7}{16}$ |

Possible positive numbers



- If, in the course of a computation, a number x is produced of the form $\pm q \times 2^m$, where m is outside the computer's permissible range, then we say that an overflow or an underflow has occurred or that x is outside the range of the computer.
- Generally, an overflow results in a fatal error (or exception), and the normal execution of the program stops.
- An underflow, however, is usually treated automatically by setting x to zero without any interruption of the program but with a warning message in most computers.
- In this specific example, any number closer to zero than $1/16$ would underflow to zero, and any number outside the range -1.75 to $+1.75$ would overflow to machine infinity.

Hole at Zero


- If, in this Example, we allow only normalized floating-point numbers, then all our numbers (with the exception of zero) have the form

$$x = \pm(0.1b_2b_3)_2 \times 2^{\pm k}$$

- This creates a phenomenon known as the **hole at zero**. The nonnegative machine numbers are now distributed as in Figure below.
- There is a relatively wide gap between zero and the smallest positive machine number, which is $(0.100)_2 \times 2^1 = \frac{1}{4}$

5-bits to represent this F-P example

- These normalized floating-point numbers can be stored in a 5-bit computer with 1 bit for sign of the number, two bits for exponent (in which 1 bit is for exponent sign) and two bits for mantissa



- All possible combinations of positive normalized F-P numbers are

$$(0.1b_2b_3)_2 \times 2^m = \left\{ \begin{array}{l} (0.100)_2 = \frac{1}{2} \\ (0.101)_2 = \frac{5}{8} \\ (0.110)_2 = \frac{3}{4} \\ (0.111)_2 = \frac{7}{8} \end{array} \right\} \times 2^{-1,0,1} = \left\{ \begin{array}{l} \frac{1}{4}, \frac{1}{2}, 1 \\ \frac{5}{16}, \frac{5}{8}, \frac{5}{4} \\ \frac{3}{8}, \frac{3}{4}, \frac{3}{2} \\ \frac{7}{16}, \frac{7}{8}, \frac{7}{4} \end{array} \right.$$

- So, a machine number in floating-point single-precision is of the form

$$(-1)^s q \times 2^m = (-1)^s \times 2^{c-1} \times (1.b_2b_3)_2$$

IEEE-754 Floating Point Standard

- Standardizes representation of floating point numbers on different computers in single and double precision.
- Standardizes representation of floating point operations on different computers.

| Precision | Bits | Sign | Exponent | Mantissa |
|-------------|------|------|----------|----------|
| Single | 32 | 1 | 8 | 23 |
| Double | 64 | 1 | 11 | 52 |
| Long Double | 80 | 1 | 15 | 64 |


IEEE-754 Format Single Precision

$$x = \pm q \times 2^m$$

- sign of q: 1 bit
- integer $|m|$: 8 bits
- number q: 23 bits

$$(-1)^s \times 2^{c-127} \times (1.f)_2$$

32 bits for single precision




Sign
(s)

Biased
Exponent (m)

Mantissa (f)

Single-Precision Floating-Point Form

- In the normalized representation of a nonzero floating-point number, the first bit in the mantissa is always 1 so that this bit does not have to be stored.
- This can be accomplished by shifting the binary point to a “1-plus” form $(1.f)_2$.
- The mantissa is the rightmost 23 bits and contains f with an understood binary point as in this Figure .
- So the mantissa (significand) actually corresponds to 24 binary digits since there is a hidden bit. (An important exception is the number ± 0 .)



Single-Precision Floating-Point Form

- The value of c in the representation of a floating-point number in single precision is restricted by the inequality

$$0 < c < (11\ 111\ 111)_2 = 255$$

- The values 0 and 255 are reserved for special cases, including ± 0 and $\pm\infty$, respectively.
- Hence, the actual exponent of the number is restricted by the inequality

$$-126 \leq c-127 \leq 127$$

- Likewise, we find that the mantissa of each nonzero number is restricted by the inequality
- $1 \leq (1.f)_2 \leq (1.111\ 111\ 111\ 111\ 111\ 111\ 111\ 11)_2 = 2-2^{-23}$

Single-Precision Floating-Point Form

- The largest number representable is therefore $(2 - 2^{-23})2^{127} \approx 2^{128} \approx 3.4 \times 10^{38}$.
- The smallest positive number is $2^{-126} \approx 1.2 \times 10^{-38}$
- The binary machine floating-point number $\varepsilon = 2^{-23}$ is called the machine epsilon when using single precision. It is the smallest positive machine number ε such that $1 + \varepsilon \neq 1$.
- Because $2^{-23} \approx 1.2 \times 10^{-7}$, we infer that in a simple computation, approximately six significant decimal digits of accuracy may be obtained in single precision. Recall that 23 bits are allocated for the mantissa.

Special Representations

- Special cases:
- 0 represented as sign bit $s = 0$ or 1 , $c = 0$ and $f = 0$ (all zeros)
- $\pm\infty$ represented as sign bit $s = 0$ or 1 , $c = 255$ and $f = 0$ (all zeros)
- NaN (division by 0) represented as $c = 255$ and $f \neq 0$

| s | c | f | Represents |
|--------|-----------|-----------|------------|
| 0 | all zeros | all zeros | 0 |
| 1 | all zeros | all zeros | -0 |
| 0 | all ones | all zeros | $+\infty$ |
| 1 | all ones | all zeros | $-\infty$ |
| 0 or 1 | all ones | non-zero | NaN |

Example 1

- Determine the single-precision machine representation of the decimal number -52.234375
- The whole part is $(52.)_{10} = (64.)_8 = (110\ 100.)_2$
- The fractional part, we have $(.234375)_{10} = (.17)_8 = (.001\ 111)_2$
- $(52.234375)_{10} = (110\ 100.001\ 111)_2 = (1.101\ 000\ 011\ 110)_2 \times 2^5$
- $(.101\ 000\ 011\ 110)_2$ is the stored mantissa.
- The exponent is $(5)_{10}$, and since $c-127 = 5$
- Hence $c = (132)_{10} = (204)_8 = (10\ 000\ 100)_2$ is the stored exponent.
- So, the single-precision machine representation
$$[1\ 10\ 000\ 100\ 101\ 000\ 011\ 110\ 000\ 000\ 000\ 00]_2 =$$
$$[1100\ 0010\ 0101\ 0000\ 1111\ 0000\ 0000\ 0000]_2 = [C250F000]_{16}$$

Example 2

- What is the single-precision representation of $(24.625)_{10}$

$$(24.)_{10} = (11000.)_2 \text{ and } (0.625)_{10} = (0.101)_2$$

$$\text{Hence, } (24.625)_{10} = (11000.101)_2 = (1.1000101)^2 = (1.1000101) \times 2^4$$

$$c - 127 = 4, \text{ so } c = 131 = (10000011)_2$$

So, the single-precision representation is

$$[0 \ 10 \ 000 \ 011 \ 100 \ 010 \ 100 \ 000 \ 000 \ 000 \ 000 \ 00]_2 =$$

$$[0100 \ 0001 \ 1100 \ 1010 \ 0000 \ 0000 \ 0000 \ 0000]_2 = [41CA0000]_{16}$$

Single Precision to Decimal

- What decimal number is represented by
01000001011111000000000000000000

Sign bit is 0, so positive number

$$c = (10000010)_2 = 130, \text{ so } m = c - 127 = 3$$

$$\text{mantissa } f = (1.111110000000000000000000)_2$$

So, number is

$$(1.111110000000000000000000)_2 \times 2^3$$

$$= (1111.11000000000000000000)_2$$

$$= 15.0 + 0.5 + 0.25 = (15.75)_{10}$$

Double Precision Floating Point Representation

- In double precision, there are 52 bits allocated for the mantissa. The double precision machine epsilon is $2^{-52} \approx 2.2 \times 10^{-16}$, so approximately 15 significant decimal digits of precision are available.
- There are 11 bits allowed for the exponent, which is biased by 1023.
- Finally, 64 bits represent the double precision number.
- The exponent represents numbers from -1022 through 1023.
- A machine number in standard double precision floating-point form corresponds to

$$(-1)^s \times 2^{c-1023} \times (1.f)_2$$

Double Precision Floating Point Representation

- The value of c in the representation of a floating-point number in double precision is restricted by the inequality

$$0 < c < (1\ 111\ 111\ 111)_2 = 2047$$

- As in single precision, the values at the ends of this interval are reserved for special cases.
- Hence, the actual exponent of the number is restricted by the inequality

$$-1022 \leq c-1023 \leq 1023$$

- We find that the mantissa of each nonzero number is restricted by the inequality
- $1 \leq (1.f)_2 \leq (1.111\ 111\ 111 \cdots 111\ 111\ 111\ 1)_2 = 2-2^{-52}$

Double Precision Floating Point Representation

- Recall that 52 bits are allocated for the mantissa.
- The largest double-precision machine number is $(2 - 2^{-52})2^{1023} \approx 2^{1024} \approx 1.8 \times 10^{308}$.
- The smallest double-precision positive machine number is $2^{-1022} \approx 2.2 \times 10^{-308}$.
- Consequently, the range for integers is from $-(2^{31} - 1)$ to $(2^{31} - 1) = 2147483647$.
- In double precision, 63 bits are used for integers giving integers in the range $-(2^{63} - 1)$ to $(2^{63} - 1)$.

Double-Precision Representation

- -52.234375 in double precision
- for the exponent $(5)_{10}$, we let $c-1023=5$, and we have $(1028)_{10} = (2004)_8 = (10\ 000\ 000\ 100)_2$, which is the stored exponent.
- Thus, the double-precision machine representation of -52.234375 is
- $[1100000001001010000011110000 \dots 00]_2 =$
 $[11000000010010100000111100000 \dots 0000]_2 = [C04A1E000000000000]_{16}$
- Class Exercise: Find Double-Precision representation of
 -285.75
 $[C071DC000000000000]_{16}$